

Lab3 Report by Enzo Zhu and Shixun WU

This report describes our work on Lab 3. Section 1, 2 describes our implementation of Part 1, and 2. Section 3 demonstrates our experimental results. Here is the link to our video

<https://drive.google.com/file/d/1tfDxuuGesUYgBNlttsX1xB672o0Gp48/view?usp=sharing>.

Enzo Zhu and Shixun Wu contribute to this project equally.

Part 1:

PCB modification

We first add two int32 variables, namely `thread_id`, `num_thread` into the PCB struct as follows:

```
// Per-process state
struct proc {
    // Lab3 addition
    int thread_id; // thread id of a process
    int num_thread; // number of thread of a process
};
```

allocproc() modification

The newly added variables are initialized in function `allocproc` as follows:

```
found:
    p->thread_id = 0;
    p->num_thread = 0;
```

New kernel function allocproc_thread()

We add a function `allocproc_thread(int pid, int tid, pagetable_t pagetable)`. In this function, we first find an unused proc struct from process array. We initialize the `thread_id`, `pagetable` with the input argument. The `thread_id` is obtained from the variable `num_thread` of parent process. The `pagetable` is same with the parent process. After that, we allocate one page of `trapframe` for this thread. Then we map the physical address of `trapframe` below other `trapframe` at the top of the process's shared virtual address using the following code `mappages(p->pagetable, TRAPFRAME - PGSIZE * (p->thread_id), PGSIZE, (uint64)(p->trapframe), PTE_R | PTE_W)`. Finally, we empty the `context` of the thread, initialize `context.ra` with `forkret`, and `context.sp` with `p->kstack + PGSIZE`;

```
struct proc*
allocproc_thread(int pid, int tid, pagetable_t pagetable)
{
    struct proc *p;
```

```

for(p = proc; p < &proc[NPROC]; p++) {
    acquire(&p->lock);
    if(p->state == UNUSED) {
        goto found;
    } else {
        release(&p->lock);
    }
}
return 0;

found:
p->pid = pid;
p->state = USED;
p->num_syscall = 0;
p->tick = 0;
p->thread_id = tid;
p->tickets = 10000;
p->pagetable = pagetable;
#ifdef STRIDE
p->stride = 1;
p->pass = 0;
#endif

// Allocate a trapframe page.
if((p->trapframe = (struct trapframe *)kalloc()) == 0){
    freeproc(p);
    release(&p->lock);
    return 0;
}

if(mappages(p->pagetable, TRAPFRAME - PGSIZE * (p->thread_id), PGSIZE,
            (uint64)(p->trapframe), PTE_R | PTE_W) < 0){
    uvmunmap(pagetable, TRAMPOLINE, 1, 0);
    freeproc(p);
    release(&p->lock);
    return 0;
}
memset(&p->context, 0, sizeof(p->context));
p->context.ra = (uint64)forkret;
p->context.sp = p->kstack + PGSIZE;

return p;
}

```

New syscall clone()

The system call `clone()` is initialized with syscall number 26. In `clone()`, we first fetch the argument `stack` with function `argaddr()`. After that, we increment the `num_thread` of caller process by one. The `thread_id` of the cloned thread is initialized with current value of `num_thread`. The `pid` and `pagetable` of the cloned thread is the same with parent process. After that, the `allocproc_thread()` is called to initialize the PCB, trapframe and context of the cloned thread. When the `allocproc_thread()` returns, the `clone()` syscall

initialized the content of `trapframe`, `sz`, file descriptor `opfile`, runnable variable `state`, and `parent` variables. The implementations of `clone()` is as follow

```
uint64 sys_clone(void)
{
    uint64 stack;
    argaddr(0, &stack);
    if(stack == 0) return -1;
    int i;
    struct proc *np;
    struct proc *p = myproc();

    acquire(&p->lock);
    p->num_thread++;

    int tid = p->num_thread;
    int npid = p->pid;
    pagetable_t npagetable = p->pagetable;
    release(&p->lock);
    // Allocate process.
    if((np = allocproc_thread(npid, tid, npagetable)) == 0){
        return -1;
    }

    np->sz = p->sz;

    // copy saved user registers.
    *(np->trapframe) = *(p->trapframe);

    // Cause fork to return 0 in the child.
    np->trapframe->a0 = 0;
    np->trapframe->sp = (uint64)stack;

    // increment reference counts on open file descriptors.
    for(i = 0; i < NOFILE; i++)
        if(p->ofile[i])
            np->ofile[i] = filedup(p->ofile[i]);
    np->cwd = idup(p->cwd);

    safestrcpy(np->name, p->name, sizeof(p->name));

    release(&np->lock);

    acquire(&wait_lock);
    np->parent = p;
    release(&wait_lock);

    acquire(&np->lock);
    np->state = RUNNABLE;
    release(&np->lock);
}
```

```

    return tid;
}

```

Modification on wait() and freeproc()

Instead of modifying `wait()`, we rewrite the kernel function `freeproc()`. We add a branch for thread and process in `freeproc()`. If the `proc *p` is a thread, namely `p->thread_id > 0`, we remove the `trapframe` from the `pagetable` using the following code `uvmunmap(p->pagetable, TRAPFRAME - (PGSIZE * p->thread_id), 1, 0);`. Unlike the process, free a thread does not free the `pagetable`.

```

static void
freeproc(struct proc *p)
{
    if(p->thread_id != 0){
        if(p->trapframe)
            uvmunmap(p->pagetable, TRAPFRAME - (PGSIZE * p->thread_id), 1, 0);
        p->trapframe = 0;
        p->pagetable = 0;
        p->thread_id = 0;
        p->sz = 0;
        p->pid = 0;
        p->parent = 0;
        p->name[0] = 0;
        p->chan = 0;
        p->killed = 0;
        p->xstate = 0;
        p->state = UNUSED;
    }
    else{
        if(p->trapframe)
            kfree((void*)p->trapframe);
        p->trapframe = 0;
        if(p->pagetable)
            proc_freepagetable(p->pagetable, p->sz);
        p->pagetable = 0;
        p->sz = 0;
        p->pid = 0;
        p->parent = 0;
        p->name[0] = 0;
        p->chan = 0;
        p->killed = 0;
        p->xstate = 0;
        p->state = UNUSED;
    }
}

```

Part 2:

thread_create() routine

The `thread_create()` function first allocate one page of user stack, and then call the `clone()` syscall. The newly cloned thread obtains a zero return value, then the `start_routine` is called. The parent process return directly.

```
int thread_create(void *(start_routine)(void*), void *arg){
    uint64* user_stack = malloc(PGSIZE) + PGSIZE;
    int ret = clone((void*) user_stack);
    if(ret < 0) return -1;
    if(ret != 0) return 0;
    start_routine(arg);
    exit(0);
}
```

spinlock routine

When a lock is initialized, the `lock-locked` is set to zero. In `lock_acquire()`, `while(__sync_lock_test_and_set(&lock->locked, 1) != 0)` is used to acquire the lock. After that `__sync_synchronize();` is called for synchronization. In `lock_release()`, `__sync_lock_release(&lock->locked);` is called after the synchronization.

```
struct lock_t {
    uint locked;
};

void
lock_init(struct lock_t* lock)
{
    lock->locked = 0;
}

// Acquire the lock.
// Loops (spins) until the lock is acquired.
void
lock_acquire(struct lock_t *lock)
{
    while(__sync_lock_test_and_set(&lock->locked, 1) != 0)
        ;
    __sync_synchronize();
}

// Release the lock.
void
lock_release(struct lock_t *lock)
{
    __sync_synchronize();
    __sync_lock_release(&lock->locked);
}
```

Part 3: Experimental Results

```
xv6 kernel is booting

hart 2 starting
hart 1 starting
init: starting sh
$ lab3_test 10 3
Round 1: thread 0 is passing the token to thread 1
Round 2: thread 1 is passing the token to thread 2
Round 3: thread 2 is passing the token to thread 0
Round 4: thread 0 is passing the token to thread 1
Round 5: thread 1 is passing the token to thread 2
Round 6: thread 2 is passing the token to thread 0
Round 7: thread 0 is passing the token to thread 1
Round 8: thread 1 is passing the token to thread 2
Round 9: thread 2 is passing the token to thread 0
Round 10: thread 0 is passing the token to thread 1
Frisbee simulation has finished, 10 rounds played in total
$ lab3_test 21 20
Round 1: thread 0 is passing the token to thread 1
Round 2: thread 1 is passing the token to thread 2
Round 3: thread 2 is passing the token to thread 3
Round 4: thread 3 is passing the token to thread 4
Round 5: thread 4 is passing the token to thread 5
Round 6: thread 5 is passing the token to thread 6
Round 7: thread 6 is passing the token to thread 7
Round 8: thread 7 is passing the token to thread 8
Round 9: thread 8 is passing the token to thread 9
Round 10: thread 9 is passing the token to thread 10
Round 11: thread 10 is passing the token to thread 11
Round 12: thread 11 is passing the token to thread 12
Round 13: thread 12 is passing the token to thread 13
Round 14: thread 13 is passing the token to thread 14
Round 15: thread 14 is passing the token to thread 15
Round 16: thread 15 is passing the token to thread 16
Round 17: thread 16 is passing the token to thread 17
Round 18: thread 17 is passing the token to thread 18
Round 19: thread 18 is passing the token to thread 19
Round 20: thread 19 is passing the token to thread 0
Round 21: thread 0 is passing the token to thread 1
Frisbee simulation has finished, 21 rounds played in total
$
```